

# **Lab Sheet 1**

**Sabaragamuwa University of Sri Lanka**

**Faculty of Computing**

**Software Engineering**

**SE6103 - Parallel and Distributed Systems**

Name : R.D.M. Perera

Reg No : 19APSE4287

Academic Period : 3<sup>rd</sup> Year 2<sup>nd</sup> Semester

Due Date : 18/11/2024

## Check Docker Version

docker --version

```
PS C:\Users\User> docker --version
Docker version 24.0.5, build ced0996
```

## Checking the Images which are already downloaded

docker images

```
PS C:\Users\User> docker images
REPOSITORY          TAG                 IMAGE ID            CREATED             SIZE
dockerapp            1.1                e251ed860845       3 weeks ago        159MB
nginx                latest             3b25b682ea82       6 weeks ago        192MB
alpine               latest             2df64a0493bc       8 months ago       7.38MB
mevin                0.0.1             2df64a0493bc       8 months ago       7.38MB
mevin                0.0.2             2df64a0493bc       8 months ago       7.38MB
<none>              <none>            c3a1088472dd       8 months ago       7.38MB
postgres             latest             b9390dd1ea18       9 months ago       431MB
dpage/pgadmin4       7.7               ce8ae42da4a0       14 months ago      545MB
postgres             15.4              68a92c148701       14 months ago      411MB
hello-world          latest             d2c94e258dc6       18 months ago      13.3kB
bde2020/hadoop-namenode latest             b638307a2119       4 years ago         1.37GB
sequenceiq/hadoop-docker 2.7.1             42efa33d1fa3       8 years ago         1.76GB
```

## Run the Docker Container

docker run -it --name hadoop-cluster -p 9870:9870 -p 8088:8088 -p 50070:50070 bde2020/hadoop-namenode:latest /bin/bash

```
PS C:\Users\User> docker run -it --name hadoop-cluster -p 9870:9870 -p 8088:8088 -p 50070:50070 bde2020/hadoop-namenode:latest /bin/bash
Configuring core
- Setting fs.defaultFS=hdfs://47102cf40a88:8020
Configuring hdfs
- Setting dfs.namenode.name.dir=file:///hadoop/dfs/name
Configuring yarn
Configuring httpfs
Configuring kms
Configuring mapred
Configuring for multihomed network
root@47102cf40a88:/# hdfs namenode -format
2024-11-18 17:05:28,420 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
```

## Configure the Hadoop file system

hdfs namenode -format

hdfs namenode &

hdfs datanode &

```
root@47102cf40a88:/# hdfs namenode -format
2024-11-18 17:05:28,420 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG: host = 47102cf40a88/172.17.0.2
STARTUP_MSG: args = [-format]
STARTUP_MSG: version = 3.2.1
STARTUP_MSG: classpath = /etc/hadoop:/opt/hadoop-3.2.1/share/hadoop/common/lib/jersey-core-1.19.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/re2j-1.1.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/log4j-1.2.17.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jackson-databind-2.9.8.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-xml-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/jetty-servlet-9.3.24.v20180605.jar:/opt/hadoop-3.2.1/share/hadoop/common/lib/animal-sniffer-annotations-1.17.jar:/opt/hadoop-3.2.1/share/hadoop
```



## Start the node manager and resource manager

yarn nodemanager &

yarn resourcemanager &

## Upload sample data to hdfs

hdfs dfs -mkdir -p /user/hadoop/input

```
root@47102cf40a88:/# hdfs dfs -mkdir -p /user/hadoop/input
2024-11-18 17:17:52,322 INFO namenode.FSEditLog: Number of transactions: 4 Total time for transactions(ms): 33 Number of transactions batched in Syncs: 0 Number of syncs: 2 SyncTimes(ms): 11
```

hdfs dfs -put \$HADOOP\_HOME/etc/hadoop/\*.xml /user/hadoop/input

```
root@47102cf40a88:/# hdfs dfs -put $HADOOP_HOME/etc/hadoop/*.xml /user/hadoop/input
2024-11-18 17:19:07,616 INFO namenode.FSEditLog: Number of transactions: 5 Total time for transactions(ms): 34 Number of transactions batched in Syncs: 0 Number of syncs: 3 SyncTimes(ms): 13
2024-11-18 17:19:07,694 INFO hdfs.StateChange: BLOCK* allocate blk_1073741825_1001, replicas=172.17.0.2:9866 for /user/hadoop/input
2024-11-18 17:19:07,745 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
2024-11-18 17:19:07,948 INFO datanode.DataNode: Receiving BP-1755187175-172.17.0.2-1731949530365:blk_1073741825_1001 src:172.17.0.2:9866
2024-11-18 17:19:08,072 INFO DataNode.clienttrace: src: /172.17.0.2:37194, dest: /172.17.0.2:9866, bytes: 8260, op: HDFS_READ, offset: 0, srvID: ac52af37-d20d-41bd-b05c-c11cfc2cddb, blockid: BP-1755187175-172.17.0.2-1731949530365:blk_1073741825_1001
```

## Run the word counter job

hadoop jar \$HADOOP\_HOME/share/hadoop/mapreduce/hadoop-mapreduce-examples-\*.jar wordcount /user/hadoop/input /user/hadoop/output

```
root@47102cf40a88:/# hadoop jar $HADOOP_HOME/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar wordcount /user/hadoop/input /user/hadoop/output
2024-11-18 17:20:32,228 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2024-11-18 17:20:32,339 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2024-11-18 17:20:32,339 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2024-11-18 17:20:32,935 INFO input.FileInputFormat: Total input files to process : 9
2024-11-18 17:20:32,990 INFO mapreduce.JobSubmitter: number of splits:9
2024-11-18 17:20:33,241 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1204723782_0001
2024-11-18 17:20:33,241 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-11-18 17:20:33,494 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2024-11-18 17:20:33,495 INFO mapreduce.Job: Running job: job_local1204723782_0001
2024-11-18 17:20:33,499 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2024-11-18 17:20:33,523 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2024-11-18 17:20:33,523 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
```

## Check the output

hdfs dfs -cat /user/hadoop/output/part-r-00000

```
root@47102cf40a88:/# hdfs dfs -cat /user/hadoop/output/part-r-00000
2024-11-18 17:21:38,347 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localHostTrusted = false, remoteHostTrusted = false
"*"      21
"AS"      9
"License");  9
"alice,bob  21
"clumping"  1
(ASF)      1
(root      1
(the       9
-->       18
-1         1
-1,        1
0.0        1
1-MAX_INT. 1
```

## Stop the Cluster

exit

docker stop hadoop-cluster

```
root@47102c440a88:/# exit
```

```
exit
```

```
PS C:\Users\User> docker stop hadoop-cluster
```

```
hadoop-cluster
```

```
PS C:\Users\User>
```